

Intrusion Detection



Two additional complications

- **NOERROR:**
 - The name exists but there is no record of that given type for that name
 - For DNSSEC, prove that there is no **ds** record
 - Says the subdomain doesn't sign with DNSSEC
- **NXDOMAIN:**
 - The name does not exist
- **NSEC** (Provable denial of existence), a record with just two fields
 - Next domain name
 - The next valid name in the domain
 - Valid types for this name
 - In a bitmap for efficiency

NSEC in action

- Name is valid so **NOERROR** but no answers
- Single **NSEC** record for **www.isc.org**:
 - No names exist between **www.isc.org** and **www-dev.isc.org**
 - **www.isc.org** only has an **A**, **AAAA**, **RRSIG**, and **NSEC** record

```
nweaver% dig +dnssec TXT www.isc.org @8.8.8.8
...
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 20430
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 0, AUTHORITY: 4, ADDITIONAL: 1
...
;; QUESTION SECTION:
;www.isc.org.           IN      TXT

;; AUTHORITY SECTION:
...
www.isc.org.          3600    IN      NSEC    www-dev.isc.org. A AAAA RRSIG NSEC
www.isc.org.          3600    IN      RRSIG   NSEC {RRSIG DATA}
```

The Use of NSEC

- Proof that a name exists but no type exists for that name
 - Critical for “This subdomain doesn’t support DNSSEC”: Return an **NSEC** record with the authority stating “There is no **DS** record”
- Proof that a name does not exist
 - It falls between the two **NSEC** names
 - Plus an **NSEC** saying “there is no wildcard”
 - Provable Denial of Existence
- Allows trivial domain enumeration
 - Attacker just starts at the beginning and walks through the NSEC records
 - Some consider this bad...

So NSEC3

- Rather than having the name, use a ***hash*** of the name
 - Hash Algorithm
 - Flags
- Iterations of the hash algorithm
- Salt (optional)
- The next name
- The RRTYPEs for this name
- Otherwise acts like NSEC, just in a different space

```
nweaver% dig +dnssec TXT org @199.19.57.1
...
;; AUTHORITY SECTION:
...
h9p7u7tr2u91d0v01js911gidnp90u3h.org. 86400 IN NSEC3 1 1 1 D399EAAB
    H9Q3IMI6H6CIJ4708DK5A3HMJLEIQ0PF NS SOA RRSIG DNSKEY NSEC3PARAM
h9p7u7tr2u91d0v01js911gidnp90u3h.org. 86400 IN RRSIG NSEC3 {RRSIG}
```

Comments on NSEC3

- It doesn't ***really*** prevent enumeration
 - You get a hash-space enumeration instead, but since people chose reasonable names...
 - An attacker can just do a brute-force attack to find out what names exist and don't exist after enumerating the hash space
- The salt is pointless!
 - Since the ***whole*** name is hashed, `foo.example.com` and `foo.example.org` will have different hashes anyway
- The only way to really prevent enumeration is to ***dynamically*** sign values
 - But that defeats the purpose of DNSSEC's offline signature generation

So what can *possibly* go wrong?

- Screwups on the authority side...
 - Too many ways to count...
 - But comcast is keeping track of it:
Follow @comcastdns on twitter
- The validator can't access DNSSEC records
- The validator can't process DNSSEC records correctly

Authority Side Screwups...

- Its quite common to screw up
- Tell your registrar you support DNSSEC when you don't
 - Took down HBO Go's launch for Comcast users and those using Google Public DNS
- Rotate your key but present old signatures
- Forget that your signatures expire

And The Recursive Resolver Must Not Be Trusted!

- Most deployments validate at the recursive resolver, not the client
 - Notably Google Public DNS and Comcast
- This provides very little practical security:
 - The recursive resolver has proven to be the biggest threat in DNS
 - And this doesn't protect you between the recursive resolver and your system
- But causes a lot of headaches
 - Comcast or Google invariably get blamed when a zone screws up
 - Fortunately this is getting less common...

DNSSEC transport

- A validating client must be able to fetch the DNSSEC related records
 - It may be through the recursive resolver
 - It may be by contacting arbitrary DNS servers on the Internet
- One of these two must work or the client ***can not validate*** DNSSEC
 - This acts to limit DNSSEC's real use:
Signing other types such as cryptographic fingerprints (e.g. DANE)

Probe the Root To Check For DNSSEC Transport

- Can the client get DNSSEC data from the Internet?
 - Probe every root with DO for:
 - DS for .com with RRSIG
 - DNSKEY for . with RRSIG
 - NSEC for an invalid TLD with RRSIG
 - Serves two purposes:
 - Some networks have one or more bad root mirrors
 - Notably one Chinese educational network has root mirrors for all but 3 that don't support DNSSEC
 - If no information can be retrieved
 - Proxy which strips out DNSSEC information and/or can't handle DO

DNSSEC Root Transport: Results We've Seen In The Wild

- Bad news at Starbucks: Hotspot gateways often proxy all DNS and can't handle DO-enabled traffic
 - And then have DNS resolvers that can't handle DNSSEC requests!
- Confirmed the Chinese educational network “Bad root mirror” problem happened
 - China had local root mirrors that didn't implement DNSSEC a few years back

Implications of “No DNSSEC at Starbucks”

- DNSSEC failure depends on the usage.
- For name->address bindings:
 - If the recursive resolver practices proper port randomization:
 - No problem. The same “attackers” who can manipulate your DNS could do anything they want at the proxy that’s controlling your DNS traffic
 - Else:
 - Problem. Network is not secure
- For name->key bindings:
 - Unless the resolver supports it directly, you are Out of Luck
 - DNSSEC information must have an alternate channel if you want to use it to transmit keys instead of just IPs

In fact, my preferred DNSSEC policy For Client Validation

- For name->address mappings
 - Any existing APIs that don't provide DNSSEC status
 - If valid: use
 - If invalid OR no complete DNSSEC chain:
 - Begin an iterative fetch with the most precise DNSSEC-validated data
 - Use the result without question
- For name->data mappings
 - An API which returns DNSSEC status
 - If valid: Use
 - If invalid: Return DNSSEC failure status
 - Up to the application

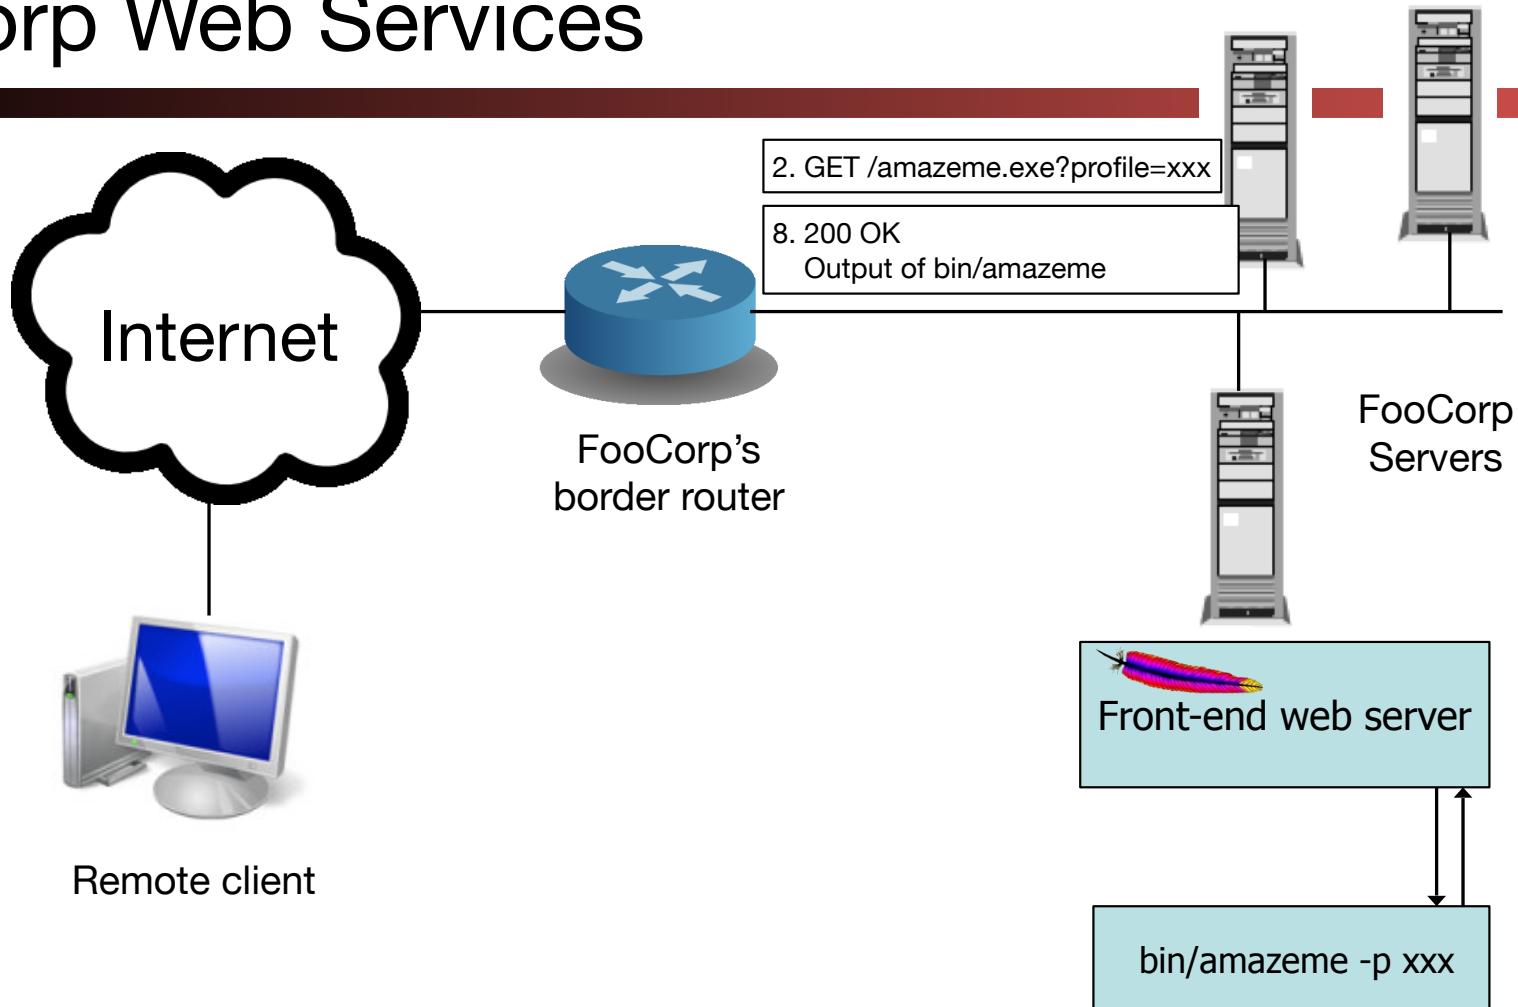
And That's The Real Thing...

- DNSSEC in all its *emm* glory.
- OPT records to say "I want DNSSEC"
- RRSIG records are certificates
- DNSKEY records hold public keys
- DS records hold key fingerprints
 - Used by the parent to tell the child's keys
- NSEC/NSEC3 records to prove that a name doesn't exist or there is no record of that type

Structure of FooCorp Web Services

Computer Science 161

Weaver



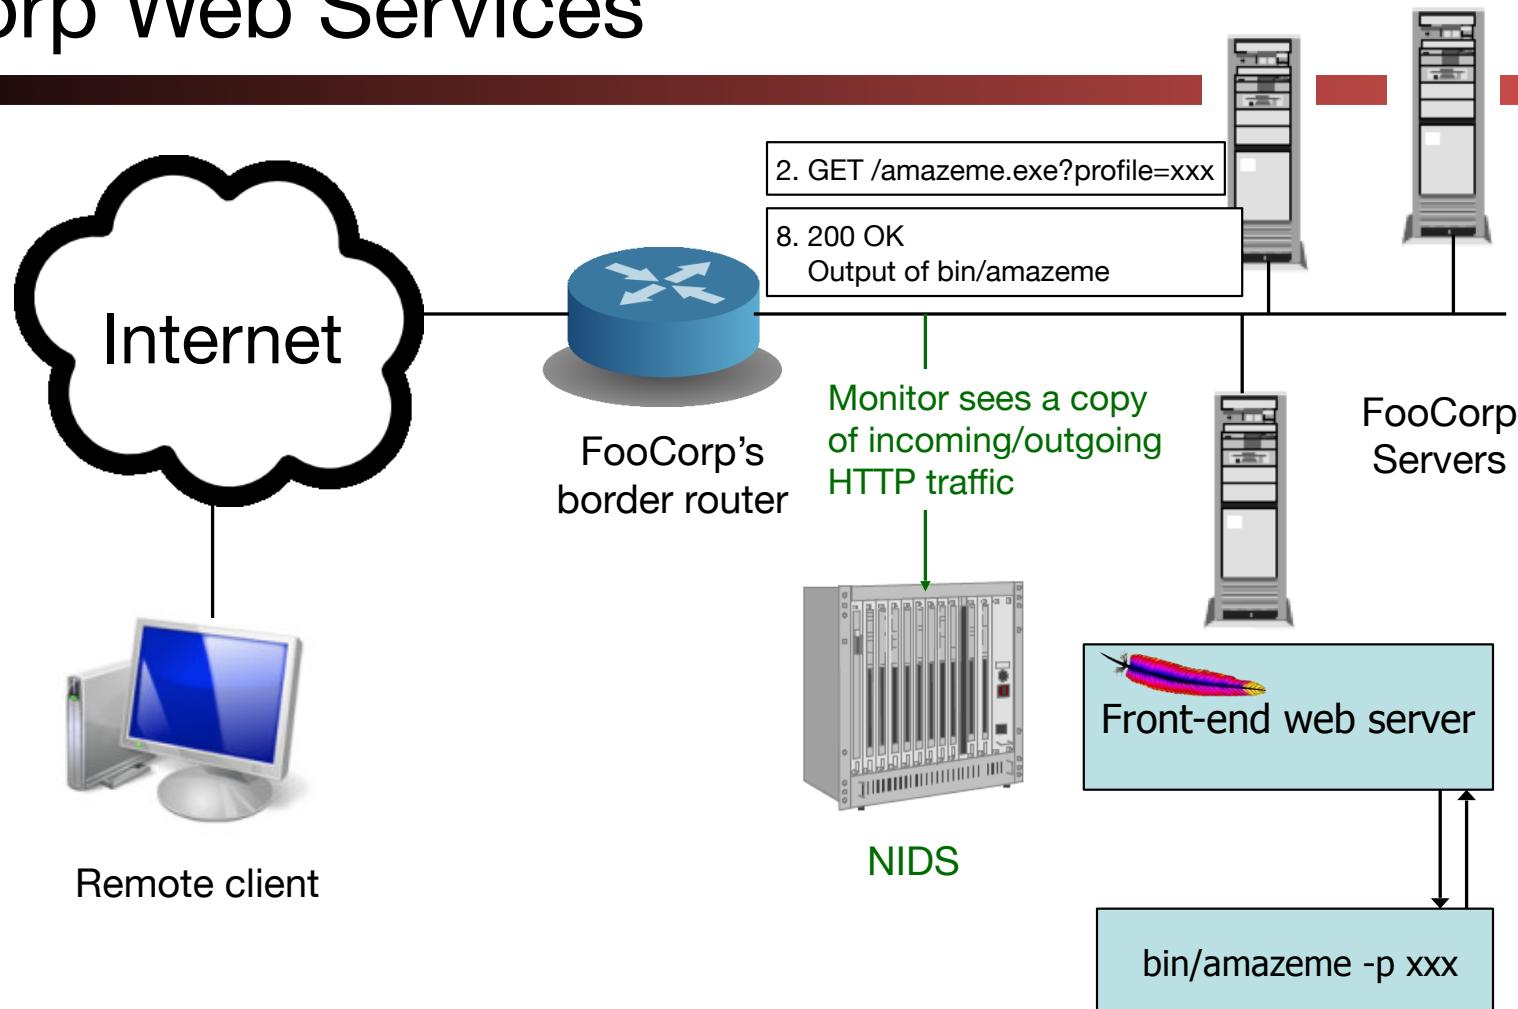
Network Intrusion Detection

- Approach #1: look at the network traffic
 - (a “NIDS”: rhymes with “kids”)
 - Scan HTTP requests
 - Look for “**/etc/passwd**” and/or “**.../.../**” in requests
 - Indicates attempts to get files that the web server shouldn't provide

Structure of FooCorp Web Services

Computer Science 161

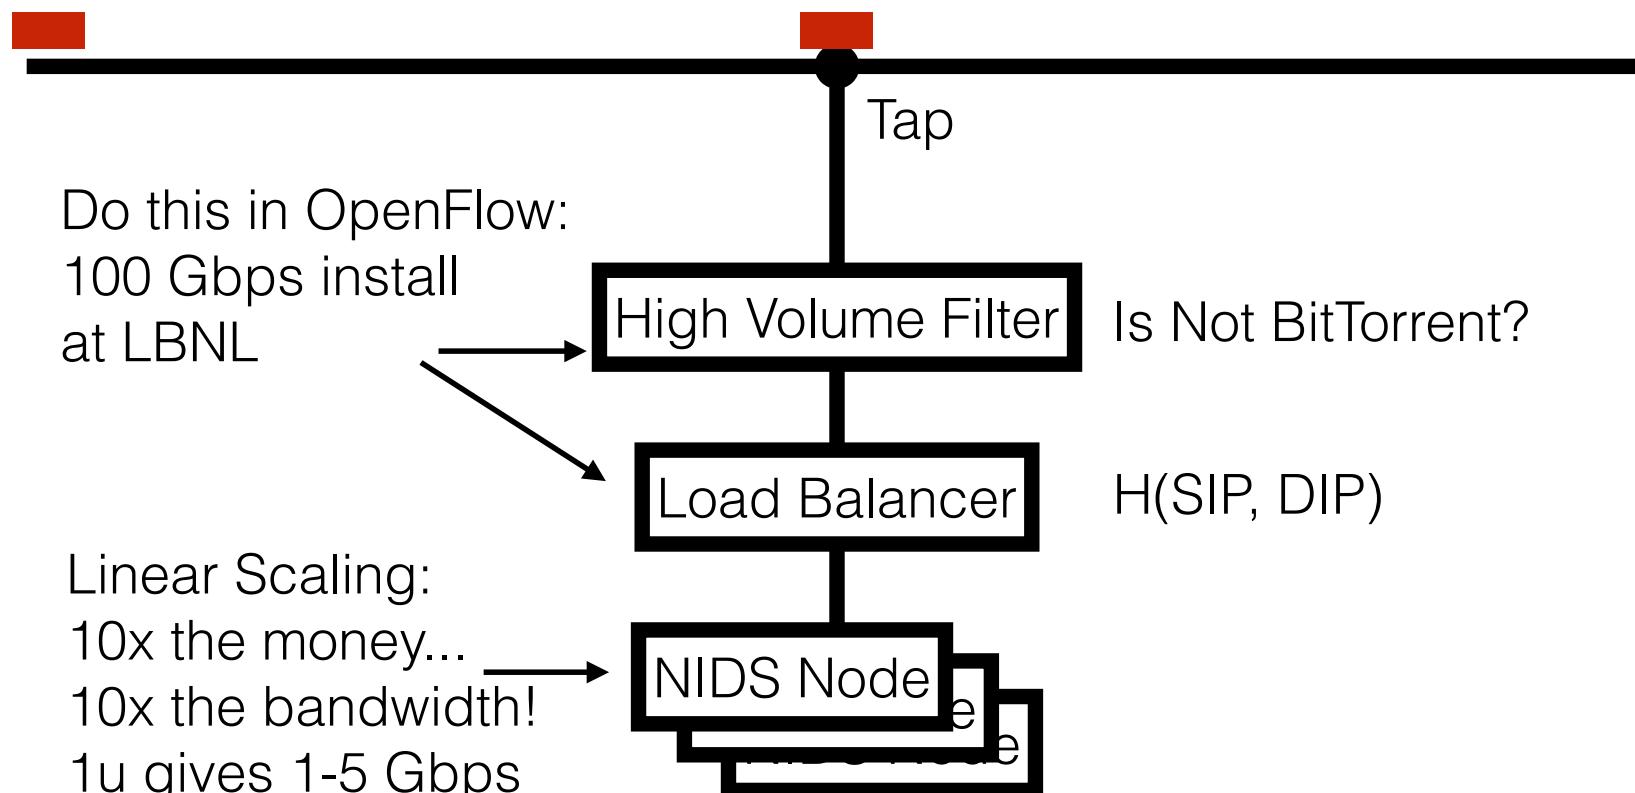
Weaver



Network Intrusion Detection

- Approach #1: look at the network traffic
 - (a “NIDS”: rhymes with “kids”)
 - Scan HTTP requests
 - Look for “/etc/passwd” and/or “.../.../”
- Pros:
 - No need to touch or trust end systems
 - Can “bolt on” security
 - Cheap: cover many systems w/ single monitor
 - Cheap: centralized management

How They Work: Scalable Network Intrusion Detection Systems



Inside the NIDS

```
GET HT TP /fu bar/ 1.1..
```

HTTP Request
URL = /fubar/
Host =

```
GET HTTP /b az/?id= 1f413 1.1...
```

HTTP Request
URL = /baz/?id=...
ID = 1f413

```
220 mail.domain.target ESMTP Sendmail...
```

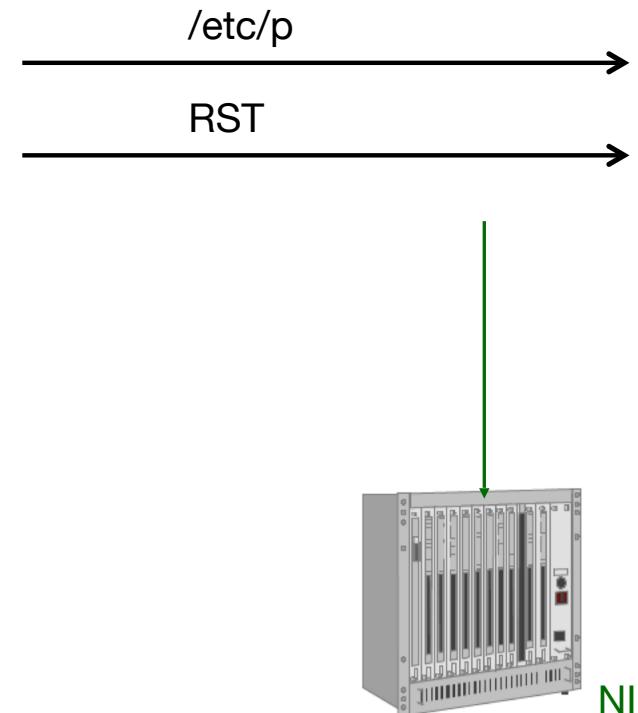
Sendmail
From = someguy@...
To = otherguy@...

Network Intrusion Detection (NIDS)

- NIDS has a table of all active connections, and maintains state for each
 - e.g., has it seen a partial match of /etc/passwd?
 - What do you do when you see a new packet not associated with any known connection?
 - Create a new connection: when NIDS starts it doesn't know what connections might be existing
 - New hotness: Network monitoring
 - Goal is not to detect attacks but just to understand everything.

Evasion

- What should NIDS do if it sees a RST packet?



- Assume RST will be received?
- Assume RST won't be received?
- Other (please specify)

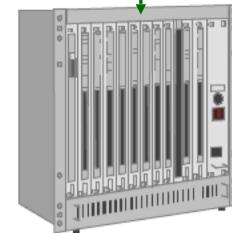
Evasion

- What should NIDS do if it sees this?

/%65%74%63/%70%61%73%73%77%64



- Alert – it's an attack
- No alert – it's all good
- Other (please specify)



NIDS

Evasion

- Evasion attacks arise when you have “double parsing”
- ***Inconsistency*** - interpreted differently between the monitor and the end system
- ***Ambiguity*** - information needed to interpret correctly is missing

Evasion Attacks (High-Level View)

- Some evasions reflect incomplete analysis
 - In our FooCorp example, hex escapes or “...///.//...” alias
 - In principle, can deal with these with implementation care (make sure we fully understand the spec)
 - Of course, in practice things inevitably fall through the cracks!
- Some are due to imperfect observability
 - For instance, if what NIDS sees doesn’t exactly match what arrives at the destination
 - EG, two copies of the “same” packet, which are actually different and with different TTLs

Network-Based Detection

- Issues:
 - Scan for “/etc/passwd”?
 - What about other sensitive files?
 - Scan for “.../.../”?
 - Sometimes seen in legit. requests (= false positive)
 - What about "%2e%2e%2f%2e%2e%2f"? (= evasion)
 - Okay, need to do full HTTP parsing
 - What about “...///...///...///”?
 - Okay, need to understand Unix filename semantics too!
 - What if it’s HTTPS and not HTTP?
 - Need access to decrypted text / session key – yuck!

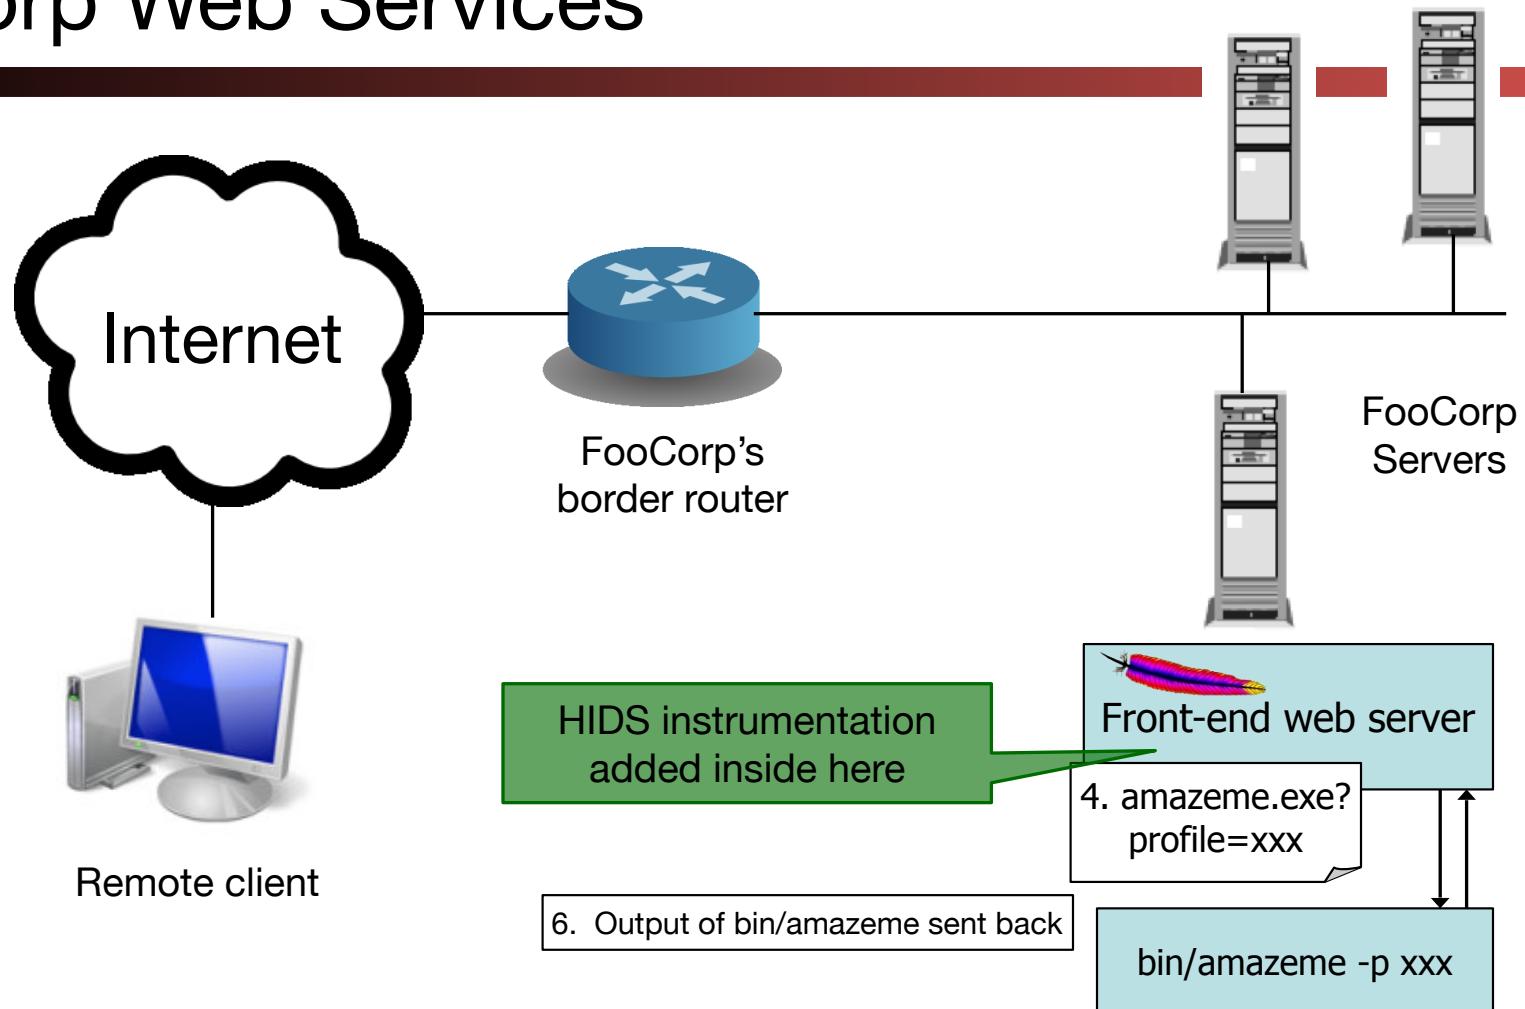
Host-based Intrusion Detection

- Approach #2: instrument the web server
 - Host-based IDS (sometimes called “HIDS”)
 - Scan ?arguments sent to back-end programs
 - Look for “/etc/passwd” and/or “..../../”

Structure of FooCorp Web Services

Computer Science 161

Weaver



Host-based Intrusion Detection

- Approach #2: instrument the web server
 - Host-based IDS (sometimes called “HIDS”)
 - Scan arguments sent to back-end programs
 - Look for “/etc/passwd” and/or “..../..”
- Pros:
 - No problems with HTTP complexities like %-escapes
 - Works for encrypted HTTPS!
- Issues:
 - Have to add code to each (possibly different) web server
 - And that effort only helps with detecting web server attacks
 - Still have to consider Unix filename semantics (“..///.//”)
 - Still have to consider other sensitive files

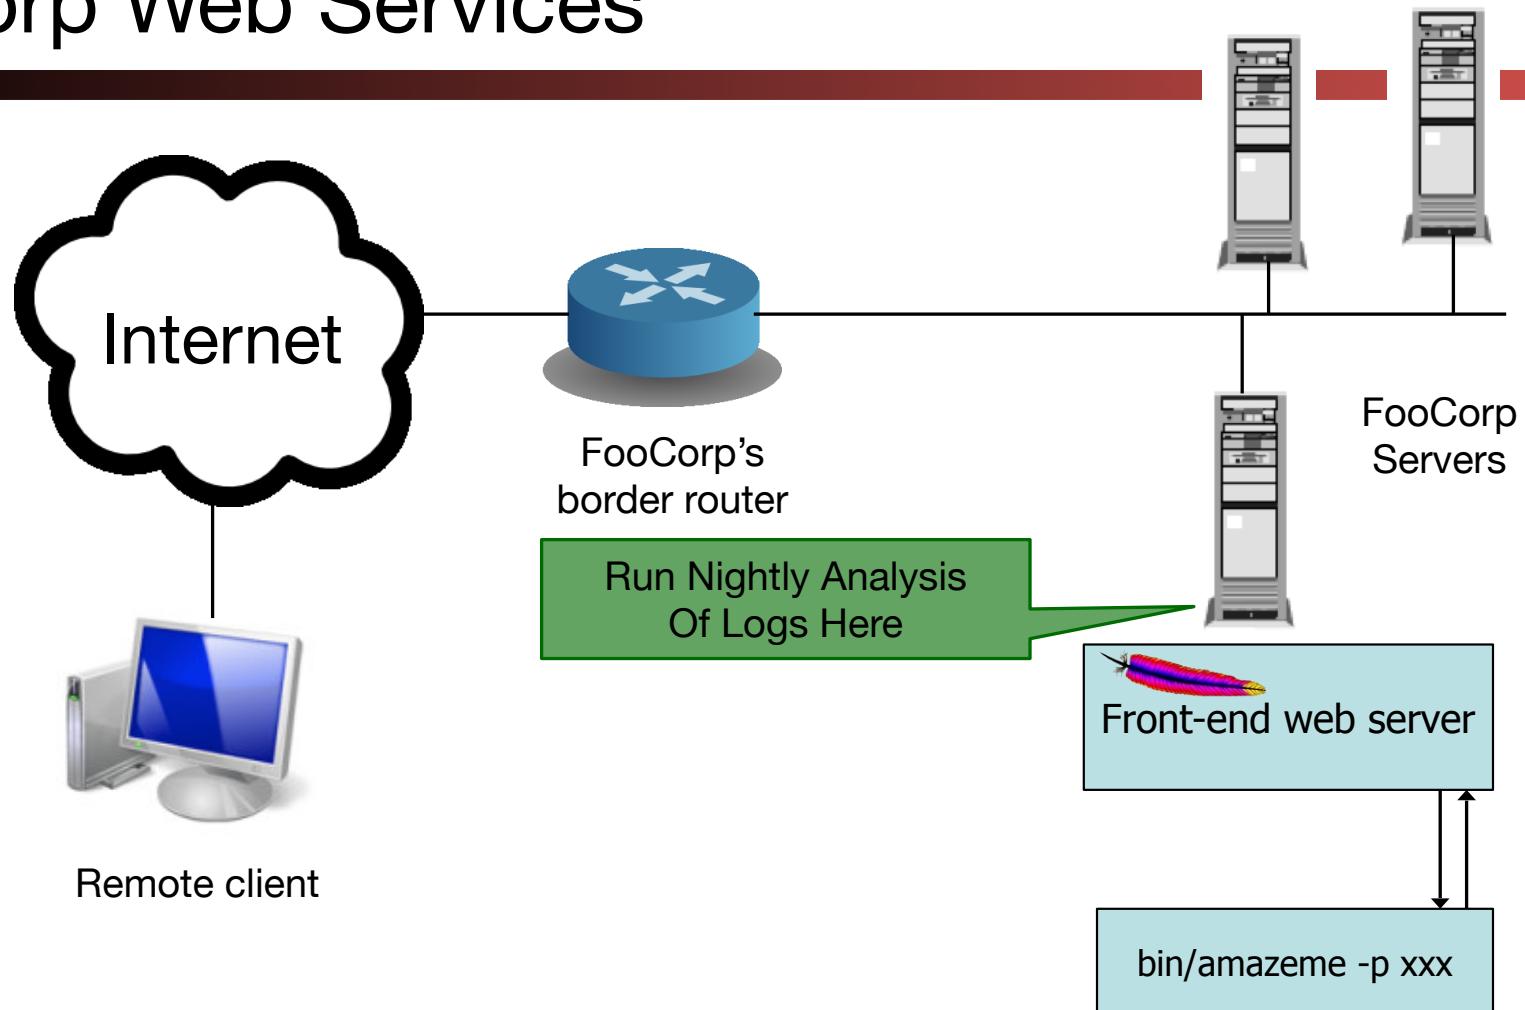
Log Analysis

- Approach #3: each night, script runs to analyze log files generated by web servers
- Again scan ?arguments sent to back-end programs

Structure of FooCorp Web Services

Computer Science 161

Weaver



Log Analysis: Aka "Log It All and let Splunk Sort It Out"

- Approach #3: each night, script runs to analyze log files generated by web servers
 - Again scan ?arguments sent to back-end programs
- Pros:
 - Cheap: web servers generally already have such logging facilities built into them
 - No problems like %-escapes, encrypted HTTPS
- Issues:
 - Again must consider filename tricks, other sensitive files
 - Can't block attacks & prevent from happening
 - Detection delayed, so attack damage may compound
 - If the attack is a compromise, then malware might be able to alter the logs before they're analyzed
 - (Not a problem for directory traversal information leak example)
 - Also can be mitigated by using a separate log server

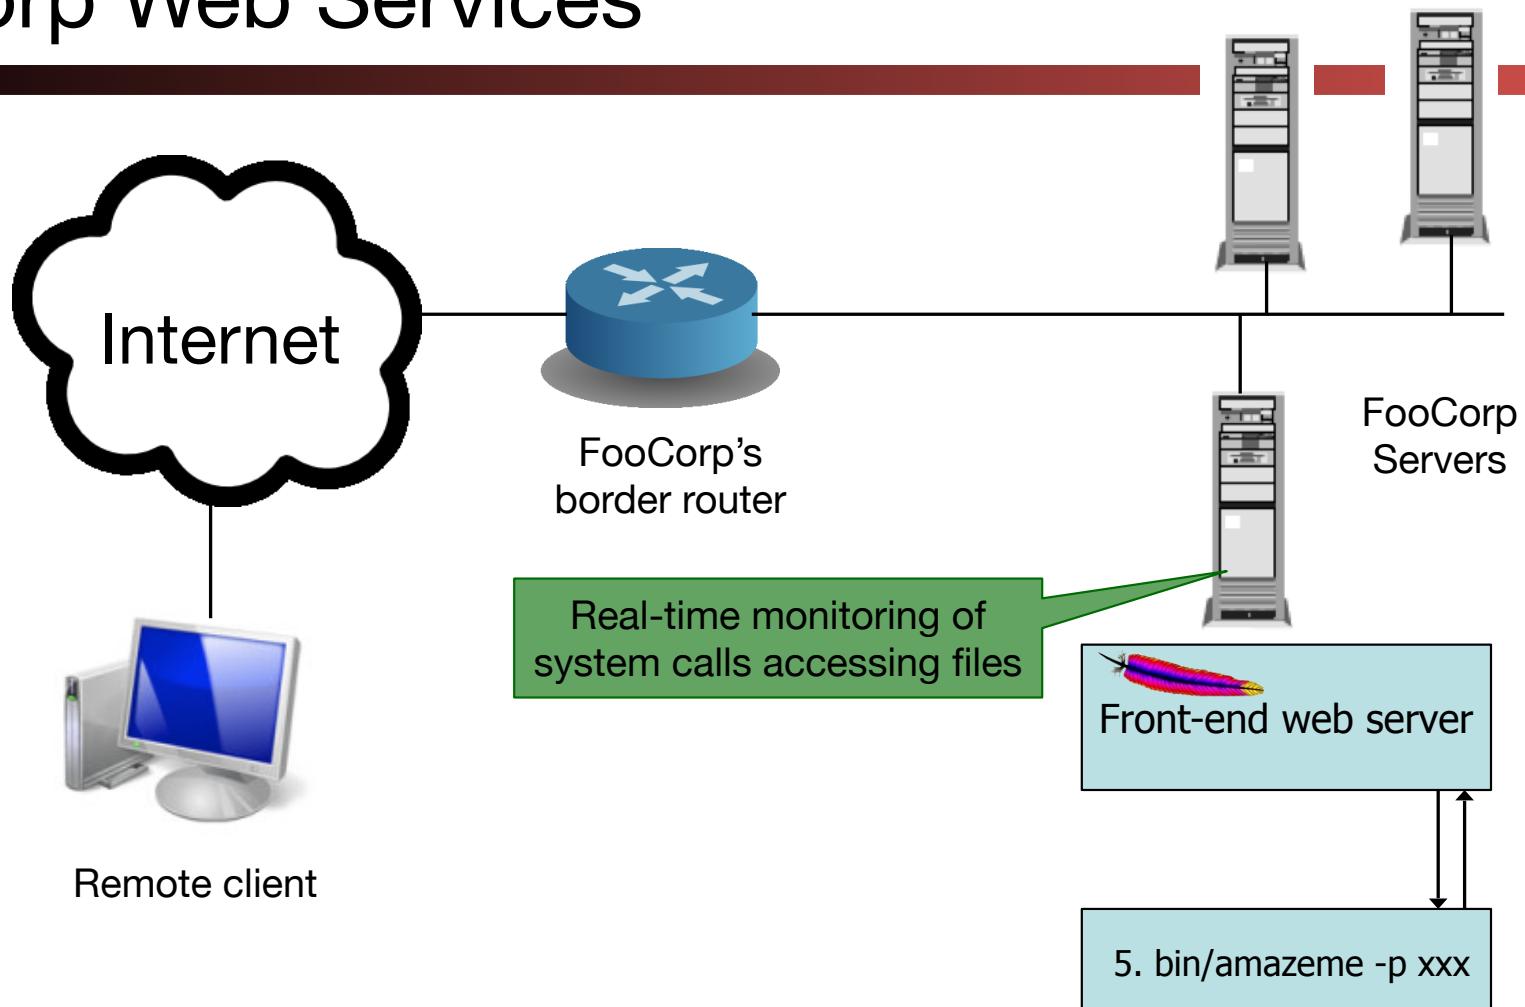
System Call Monitoring (HIDS)

- Approach #4: monitor system call activity of backend processes
- Look for access to **/etc/passwd**

Structure of FooCorp Web Services

Computer Science 161

Weaver



System Call Monitoring (HIDS)

- Approach #4: monitor system call activity of backend processes
 - Look for access to /etc/passwd
- Pros:
 - No issues with any HTTP complexities
 - May avoid issues with filename tricks
 - Attack only leads to an “alert” if attack succeeded
 - Sensitive file was indeed accessed
- Issues:
 - Maybe other processes make legit accesses to the sensitive files (false positives)
 - Maybe we’d like to detect attempts even if they fail?
 - “situational awareness”

Detection Accuracy

- Two types of detector errors:
 - False positive (FP): alerting about a problem when in fact there was no problem
 - False negative (FN): failing to alert about a problem when in fact there was a problem
- Detector accuracy is often assessed in terms of rates at which these occur:
 - Define I to be the event of an instance of intrusive behavior occurring (something we want to detect)
 - Define A to be the event of detector generating alarm
- Define:
 - False positive rate = $P[A|\neg I]$
 - False negative rate = $P[\neg A| I]$

Perfect Detection

- Is it possible to build a detector for our example with a false negative rate of 0%?
- Algorithm to detect bad URLs with 0% FN rate:

```
void my_detector_that_never_misses(char *URL)
{
    printf("yep, it's an attack!\n");
}
```

- In fact, it works for detecting any bad activity with no false negatives! Woo-hoo!
- Wow, so what about a detector for bad URLs that has NO FALSE POSITIVES?!
- `printf("nope, not an attack\n");`

Detection Tradeoffs

- The art of a good detector is achieving an effective balance between FPs and FNs
- Suppose our detector has an FP rate of 0.1% and an FN rate of 2%. Is it good enough? Which is better, a very low FP rate or a very low FN rate?
 - Depends on the cost of each type of error ...
 - E.g., FP might lead to paging a duty officer and consuming hour of their time; FN might lead to \$10K cleaning up compromised system that was missed
 - ... but also critically depends on the rate at which actual attacks occur in your environment

Base Rate Fallacy

- Suppose our detector has a FP rate of 0.1% (!) and a FN rate of 2% (not bad!)
- Scenario #1: our server receives 1,000 URLs/day, and 5 of them are attacks
 - Expected # FPs each day = $0.1\% * 995 \approx 1$
 - Expected # FNs each day = $2\% * 5 = 0.1 < 1/\text{week}$
 - Pretty good!
- Scenario #2: our server receives 10,000,000 URLs/day, and 5 of them are attacks
 - Expected # FPs each day $\approx 10,000$:-(
 - Nothing changed about the detector; only our environment changed
 - Accurate detection very challenging when base rate of activity we want to detect is quite low

Composing Detectors: There Is No Free Lunch

- "Hey, what if we take two (bad) detectors and combine them?"
 - Can we turn that into a good detector?
 - Note: Assumes the detectors are independent
- Parallel composition: Either detector triggers an alert
 - Reduces false negative rate (either one alerts works)
 - ***Increases*** false positive rate!
- Series composition: both detectors must trigger for an alert
 - Reduces false positive rate (since both must false positive)
 - ***Increases*** false negative rate!

Styles of Detection: Signature-Based

- Idea: look for activity that matches the structure of a known attack
- Example (from the freeware Snort NIDS):

```
alert tcp $EXTERNAL_NET any -> $HOME_NET 139
flow:to_server,established
content:"|eb2f 5feb 4a5e 89fb 893e 89f2|"
msg:"EXPLOIT x86 linux samba overflow"
reference:bugtraq,1816
reference:cve,CVE-1999-0811
classtype:attempted-admin
```

- Can be at different semantic layers
e.g.: IP/TCP header fields; packet payload; URLs

Signature-Based Detection

- E.g. for FooCorp, search for “`../..`” or “`/etc/passwd`”
- What’s nice about this approach?
 - Conceptually simple
 - Takes care of known attacks (of which there are zillions)
 - Easy to share signatures, build up libraries
- What’s problematic about this approach?
 - Blind to novel attacks
 - Might even miss variants of known attacks (“`.....`”)
 - Of which there are zillions
 - Simpler versions look at low-level syntax, not semantics
 - Can lead to weak power (either misses variants, or generates lots of false positives)

Vulnerability Signatures

- Idea: don't match on known attacks, match on known problems
- Example (also from Snort):

```
alert tcp $EXTERNAL_NET any -> $HTTP_SERVERS 80
uricontent: ".ida?"; nocase; dsize: > 239; flags:A+
msg:"Web-IIS ISAPI .ida attempt"
reference:bugtraq,1816
reference:cve,CAN-2000-0071
classtype:attempted-admin
```

- That is, match URIs that invoke `*.ida?*`, have more than 239 bytes of payload, and have ACK set (maybe others too)
- This example detects any* attempt to exploit a particular buffer overflow in IIS web servers
 - Used by the “Code Red” worm
 - (Note, signature is not quite complete: also worked for `*.idb?*`)

Asside: Why The Covid Vaccines Are So Good!

- The COVID vaccines are basically training your body to respond to a **vulnerability** signature!
- Not recognize any random part of the virus...
 - But the key the virus uses to invade cells!
- So if the virus mutates to avoid detection...
 - It is likely to also be less effective at invading your cells
- Plus a "head start" seems to be enough
 - These vaccines are basically 100% effective at turning COVID into a Common Cold even if you do get infected

Styles of Detection: Anomaly-Based

- Idea: attacks look peculiar.
- High-level approach: develop a model of normal behavior (say based on analyzing historical logs). Flag activity that deviates from it.
- FooCorp example: maybe look at distribution of characters in URL parameters, learn that some are rare and/or don't occur repeatedly
 - If we happen to learn that '.'s have this property, then could detect the attack even without knowing it exists
 - Big benefit: potential detection of a wide range of attacks, including novel ones

Anomaly Detection Problems

- Can fail to detect known attacks
- Can fail to detect novel attacks, if don't happen to look peculiar along measured dimension
- What happens if the historical data you train on includes attacks?
- Base Rate Fallacy particularly acute: if prevalence of attacks is low, then you're more often going to see benign outliers
 - High FP rate
 - OR: require such a stringent deviation from “normal” that most attacks are missed (high FN rate)
- Proves great subject for academic papers but not generally used

Specification-Based Detection

- Idea: don't learn what's normal; specify what's allowed
- FooCorp example: decide that all URL parameters sent to foocorp.com servers must have at most one '/' in them
 - Flag any arriving param with > 1 slash as an attack
- What's nice about this approach?
 - Can detect novel attacks
 - Can have low false positives
 - If FooCorp audits its web pages to make sure they comply
- What's problematic about this approach?
 - Expensive: lots of labor to derive specifications
 - And keep them up to date as things change ("churn")

Styles of Detection: Behavioral

- Idea: don't look for attacks, look for evidence of compromise
- FooCorp example: inspect all output web traffic for any lines that match a `passwd` file
- Example for monitoring user shell keystrokes:
`unset HISTFILE`
- Example for catching code injection: look at sequences of system calls, flag any that prior analysis of a given program shows it can't generate
 - E.g., observe process executing `read()`, `open()`, `write()`, `fork()`, `exec()` ...
 - ... but there's no code path in the (original) program that calls those in exactly that order!

Behavioral-Based Detection

- What's nice about this approach?
 - Can detect a wide range of novel attacks
 - Can have low false positives
 - Depending on degree to which behavior is distinctive
 - E.g., for system call profiling: no false positives!
 - Can be cheap to implement
 - E.g., system call profiling can be mechanized
- What's problematic about this approach?
 - Post facto detection: discovers that you definitely have a problem, w/ no opportunity to prevent it
 - Brittle: for some behaviors, attacker can maybe avoid it
 - Easy enough to not type “`unset HISTFILE`”
 - How could they evade system call profiling?
 - Mimicry: adapt injected code to comply w/ allowed call sequences (and can be automated!)

Summary of Evasion Issues

- Evasions arise from uncertainty (or incompleteness) because detector must infer behavior/processing it can't directly observe
 - A general problem any time detection separate from potential target
- One general strategy: impose canonical form (“normalize”)
 - E.g., rewrite URLs to expand/remove hex escapes
 - E.g., enforce blog comments to only have certain HTML tags
- Another strategy: analyze all possible interpretations rather than assuming one
 - E.g., analyze raw URL, hex-escaped URL, doubly-escaped URL ...
- Another strategy: Flag potential evasions
 - So the presence of an ambiguity is at least noted
- Another strategy: fix the basic observation problem
 - E.g., monitor directly at end systems

Inside a Modern HIDS (“AV”)

- URL/Web access blocking:
 - Prevent users from going to known bad locations
- Protocol scanning of network traffic (esp. HTTP)
 - Detect & block known attacks
 - Detect & block known malware communication
- Payload scanning
 - Detect & block known malware
 - (Auto-update of signatures for these)
- Cloud queries regarding reputation
 - Who else has run this executable and with what results?
 - What's known about the remote host / domain / URL?

Inside a Modern HIDS

- **Sandbox execution**
 - Run selected executables in constrained/monitored environment
 - Analyze:
 - System calls
 - Changes to files / registry
 - Self-modifying code (polymorphism/metamorphism)
- **File scanning**
 - Look for malware that installs itself on disk
- **Memory scanning**
 - Look for malware that never appears on disk
- **Runtime analysis**
 - Apply heuristics/signatures to execution behavior

Inside a Modern NIDS

- Deployment inside network as well as at border
 - Greater visibility, including tracking of user identity
- Full protocol analysis
 - Including extraction of complex embedded objects
 - In some systems, 100s of known protocols
- Signature analysis (also behavioral)
 - Known attacks, malware communication, blacklisted hosts/domains
 - Known malicious payloads
 - Sequences/patterns of activity
- Shadow execution (e.g., Flash, PDF programs)
- Extensive logging (in support of forensics)
- Auto-update of signatures, blacklists

NIDS vs. HIDS

- NIDS benefits:
 - Can cover a lot of systems with single deployment
 - Much simpler management
 - Easy to “bolt on” / no need to touch end systems
 - Doesn’t consume production resources on end systems
 - Harder for an attacker to subvert / less to trust
- HIDS benefits:
 - Can have direct access to semantics of activity
 - Better positioned to block (prevent) attacks
 - Harder to evade
 - Can protect against non-network threats
 - Visibility into encrypted activity
 - Performance scales much more readily (no chokepoint)
 - No issues with “dropped” packets

Key Concepts for Detection

- Signature-based vs anomaly detection
(blacklisting vs whitelisting)
- Evasion attacks
- Evaluation metrics: False positive rate, false negative rate
- Base rate problem

Detection vs. Blocking

- If we can detect attacks, how about blocking them?
- Issues:
 - Not a possibility for retrospective analysis (e.g., nightly job that looks at logs)
 - Quite hard for detector that's not in the data path
 - E.g. How can NIDS that passively monitors traffic block attacks?
 - Change firewall rules dynamically; forge RST packets
 - And still there's a race regarding what attacker does before block
 - False positives get more expensive
 - You don't just bug an operator, you damage production activity
- Today's technology/products pretty much all offer blocking
 - Intrusion prevention systems (IPS - “eye-pe-ess”)

Can We Build An IPS That Blocks All Attacks?



The Ultimately Secure DEEP PACKET INSPECTION AND APPLICATION SECURITY SYSTEM
Featuring signature-less anomaly detection and blocking technology with application awareness and layer-7 state tracking!!!

Now available in Petabyte-capable appliance form factor!*

(Formerly: The Ultimately Secure INTRUSION PREVENTION SYSTEM
Featuring signature-less anomaly detection and blocking technology!!)

An Alternative Paradigm

- Idea: rather than detect attacks, launch them yourself!
- Vulnerability scanning: use a tool to probe your own systems with a wide range of attacks, fix any that succeed
- Pros?
 - Accurate: if your scanning tool is good, it finds real problems
 - Proactive: can prevent future misuse
 - Intelligence: can ignore IDS alarms that you know can't succeed
- Issues?
 - Can take a lot of work
 - Not so helpful for systems you can't modify
 - Dangerous for disruptive attacks
 - And you might not know which these are ...
- In practice, this approach is prudent and widely used today
 - Good complement to also running an IDS

Styles of Detection: Honeypots

- Idea: deploy a sacrificial system that has no operational purpose
- Any access is by definition not authorized ...
- ... and thus an intruder
 - (or some sort of mistake)
- Provides opportunity to:
 - Identify intruders
 - Study what they're up to
 - Divert them from legitimate targets

Honeypots

- Real-world example: some hospitals enter fake records with celebrity names ...
 - ... to entrap staff who don't respect confidentiality
- What's nice about this approach?
 - Can detect all sorts of new threats
- What's problematic about this approach?
 - Can be difficult to lure the attacker
 - Can be a lot of work to build a convincing environment
 - Note: both of these issues matter less when deploying honeypots for automated attacks
 - Because these have more predictable targeting & env. needs
 - E.g. "spamtraps": fake email addresses to catching spambots
- A great honeypot: An unsecured Bitcoin wallet...
 - When your bitcoins get stolen, you know you got compromised!

Forensics

- Vital complement to detecting attacks: figuring out what happened in wake of successful attack
- Doing so requires access to rich/extensive logs
 - Plus tools for analyzing/understanding them
- It also entails looking for patterns and understanding the implications of structure seen in activity
 - An iterative process (“peeling the onion”)

Other Attacks on IDSs

- DoS: exhaust its memory
 - IDS has to track ongoing activity
 - Attacker generates lots of different forms of activity, consumes all of its memory
 - E.g., spoof zillions of distinct TCP SYNs ...
 - ... so IDS must hold zillions of connection records
- DoS: exhaust its processing
 - One sneaky form: algorithmic complexity attacks
 - E.g., if IDS uses a predictable hash function to manage connection records ...
 - ... then generate series of hash collisions
- Code injection (!)
 - After all, NIDS analyzers take as input network traffic under attacker's control ...
 - One of the CS194 projects will be on this topic...

And, of course, our monitors have bugs...

Computer Science 161 Weaver

Riverbed Technology WinPcap IPv4 ✓ IPv6 ✗

WIRESHARK

the world's foremost network protocol analyzer

Wireshark | Get Help | Develop Google™ Custom Search Search



Security Advisories

The following Wireshark releases fix serious security vulnerabilities. If you are running a vulnerable version of Wireshark you should consider upgrading.

- [wnpa-sec-2013-09](#): NTLMSSP dissector overflow, fixed in 1.8.5, 1.6.13
- [wnpa-sec-2013-08](#): Wireshark dissection engine crash, fixed in 1.8.5, 1.6.13
- [wnpa-sec-2013-07](#): DCP-ETSI dissector crash, fixed in 1.8.5, 1.6.13
- [wnpa-sec-2013-06](#): ROHC dissector crash, fixed in 1.8.5
- [wnpa-sec-2013-05](#): DTLS dissector crash, fixed in 1.8.5, 1.6.13
- [wnpa-sec-2013-04](#): MS-MMC dissector crash, fixed in 1.8.5, 1.6.13
- [wnpa-sec-2013-03](#): DTN dissector crash, fixed in 1.8.5, 1.6.13
- [wnpa-sec-2013-02](#): CLNP dissector crash, fixed in 1.8.5, 1.6.13